

DEEP LEARNING ASSIGNMENT NO. 5

Name : Manjiri Netankar

GROUP MEMBERS :

SUPRIYA MASKAR(202201040049)

SAPNA DAHIKAMBLE(202201070065)

MANJIRI NETANKAR(202201040206)

Colab Link :

<https://colab.research.google.com/drive/1cR8B567aUxFddGNSmxxUPowiKweMNKHi?usp=sharing>

GitHub Link : https://github.com/supriyamaskar/LSTM_deepLearning

Experiment 5.1:

Objective: To forecast future values of a univariate time series using LSTM-based models.

```
# 1. INSTALL REQUIRED LIBRARIES
!pip install -q pandas numpy scikit-learn matplotlib tensorflow

# 2. LOAD THE EXTRACTED CSV FILE
import pandas as pd

# Make sure this path is correct based on where your CSV is extracted
df = pd.read_csv('day_wise.csv')

# 3. PREPROCESSING
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error

df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)
df = df[['New cases']]

scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(df)

def create_sequences(data, seq_length=14):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i + seq_length])
        y.append(data[i + seq_length])
    return np.array(X), np.array(y)

seq_length = 14
X, y = create_sequences(scaled_data, seq_length)

split = int(len(X) * 0.8)
X_train, X_test = X[:split], X[split:]
```

```

y_train, y_test = y[:split], y[split:]

X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))

# 4. BUILD AND TRAIN LSTM MODEL
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense

model = Sequential()
model.add(LSTM(50, activation='relu', input_shape=(seq_length, 1)))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error',
metrics=['mae'])

history = model.fit(X_train, y_train, epochs=50, batch_size=8,
validation_split=0.1, verbose=1)

# 5. PREDICT AND EVALUATE
y_pred = model.predict(X_test)
y_pred_inv = scaler.inverse_transform(y_pred)
y_test_inv = scaler.inverse_transform(y_test)

rmse = np.sqrt(mean_squared_error(y_test_inv, y_pred_inv))
mae = mean_absolute_error(y_test_inv, y_pred_inv)

print(f"RMSE: {rmse:.2f}")
print(f"MAE : {mae:.2f}")

# 6. PLOT PREDICTION VS ACTUAL
plt.figure(figsize=(12,6))
plt.plot(y_test_inv, label='Actual', marker='o')
plt.plot(y_pred_inv, label='Predicted', marker='x')
plt.title('Prediction vs Actual - New COVID-19 Cases')
plt.xlabel('Days')
plt.ylabel('New Cases')
plt.legend()
plt.grid(True)
plt.show()

/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/
rnn.py:200: UserWarning: Do not pass an `input_shape` / `input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

Epoch 1/50
16/16 _____ 4s 45ms/step - loss: 0.0540 - mae: 0.1807 -
val_loss: 0.1548 - val_mae: 0.3901
Epoch 2/50

```

```
16/16 _____ 1s 26ms/step - loss: 0.0287 - mae: 0.1439 -  
val_loss: 0.0692 - val_mae: 0.2584  
Epoch 3/50  
16/16 _____ 0s 24ms/step - loss: 0.0111 - mae: 0.0950 -  
val_loss: 0.0113 - val_mae: 0.0946  
Epoch 4/50  
16/16 _____ 1s 63ms/step - loss: 0.0068 - mae: 0.0652 -  
val_loss: 0.0102 - val_mae: 0.0887  
Epoch 5/50  
16/16 _____ 1s 53ms/step - loss: 0.0050 - mae: 0.0587 -  
val_loss: 0.0055 - val_mae: 0.0601  
Epoch 6/50  
16/16 _____ 1s 70ms/step - loss: 0.0036 - mae: 0.0512 -  
val_loss: 0.0024 - val_mae: 0.0333  
Epoch 7/50  
16/16 _____ 1s 56ms/step - loss: 0.0025 - mae: 0.0446 -  
val_loss: 0.0028 - val_mae: 0.0376  
Epoch 8/50  
16/16 _____ 1s 43ms/step - loss: 0.0020 - mae: 0.0391 -  
val_loss: 0.0040 - val_mae: 0.0513  
Epoch 9/50  
16/16 _____ 0s 20ms/step - loss: 0.0019 - mae: 0.0357 -  
val_loss: 0.0025 - val_mae: 0.0339  
Epoch 10/50  
16/16 _____ 0s 18ms/step - loss: 0.0022 - mae: 0.0374 -  
val_loss: 0.0044 - val_mae: 0.0553  
Epoch 11/50  
16/16 _____ 0s 16ms/step - loss: 0.0014 - mae: 0.0294 -  
val_loss: 0.0028 - val_mae: 0.0382  
Epoch 12/50  
16/16 _____ 0s 16ms/step - loss: 0.0014 - mae: 0.0298 -  
val_loss: 0.0025 - val_mae: 0.0341  
Epoch 13/50  
16/16 _____ 0s 18ms/step - loss: 0.0015 - mae: 0.0312 -  
val_loss: 0.0025 - val_mae: 0.0337  
Epoch 14/50  
16/16 _____ 0s 16ms/step - loss: 0.0012 - mae: 0.0278 -  
val_loss: 0.0025 - val_mae: 0.0340  
Epoch 15/50  
16/16 _____ 0s 15ms/step - loss: 0.0014 - mae: 0.0288 -  
val_loss: 0.0026 - val_mae: 0.0356  
Epoch 16/50  
16/16 _____ 0s 15ms/step - loss: 0.0013 - mae: 0.0271 -  
val_loss: 0.0024 - val_mae: 0.0329  
Epoch 17/50  
16/16 _____ 0s 15ms/step - loss: 0.0013 - mae: 0.0294 -  
val_loss: 0.0027 - val_mae: 0.0371  
Epoch 18/50  
16/16 _____ 0s 15ms/step - loss: 0.0010 - mae: 0.0265 -
```

val_loss: 0.0025 - val_mae: 0.0340
Epoch 19/50
16/16 _____ 0s 17ms/step - loss: 0.0013 - mae: 0.0300 -
val_loss: 0.0024 - val_mae: 0.0338
Epoch 20/50
16/16 _____ 0s 16ms/step - loss: 0.0011 - mae: 0.0252 -
val_loss: 0.0029 - val_mae: 0.0397
Epoch 21/50
16/16 _____ 0s 16ms/step - loss: 0.0012 - mae: 0.0269 -
val_loss: 0.0028 - val_mae: 0.0390
Epoch 22/50
16/16 _____ 0s 15ms/step - loss: 8.5924e-04 - mae:
0.0226 - val_loss: 0.0033 - val_mae: 0.0436
Epoch 23/50
16/16 _____ 0s 14ms/step - loss: 0.0012 - mae: 0.0269 -
val_loss: 0.0030 - val_mae: 0.0404
Epoch 24/50
16/16 _____ 0s 18ms/step - loss: 0.0010 - mae: 0.0260 -
val_loss: 0.0036 - val_mae: 0.0463
Epoch 25/50
16/16 _____ 0s 15ms/step - loss: 9.0056e-04 - mae:
0.0232 - val_loss: 0.0028 - val_mae: 0.0386
Epoch 26/50
16/16 _____ 0s 16ms/step - loss: 9.4760e-04 - mae:
0.0233 - val_loss: 0.0034 - val_mae: 0.0447
Epoch 27/50
16/16 _____ 0s 15ms/step - loss: 0.0011 - mae: 0.0273 -
val_loss: 0.0032 - val_mae: 0.0433
Epoch 28/50
16/16 _____ 0s 16ms/step - loss: 8.2738e-04 - mae:
0.0222 - val_loss: 0.0028 - val_mae: 0.0381
Epoch 29/50
16/16 _____ 0s 20ms/step - loss: 0.0011 - mae: 0.0255 -
val_loss: 0.0028 - val_mae: 0.0380
Epoch 30/50
16/16 _____ 0s 18ms/step - loss: 9.2817e-04 - mae:
0.0224 - val_loss: 0.0037 - val_mae: 0.0470
Epoch 31/50
16/16 _____ 1s 16ms/step - loss: 0.0011 - mae: 0.0255 -
val_loss: 0.0030 - val_mae: 0.0406
Epoch 32/50
16/16 _____ 0s 18ms/step - loss: 9.8017e-04 - mae:
0.0241 - val_loss: 0.0035 - val_mae: 0.0452
Epoch 33/50
16/16 _____ 0s 17ms/step - loss: 0.0012 - mae: 0.0252 -
val_loss: 0.0037 - val_mae: 0.0470
Epoch 34/50
16/16 _____ 1s 16ms/step - loss: 0.0011 - mae: 0.0248 -
val_loss: 0.0035 - val_mae: 0.0460

Epoch 35/50
16/16 _____ 0s 20ms/step - loss: 9.0387e-04 - mae: 0.0234 - val_loss: 0.0037 - val_mae: 0.0471

Epoch 36/50
16/16 _____ 1s 21ms/step - loss: 8.0381e-04 - mae: 0.0217 - val_loss: 0.0034 - val_mae: 0.0447

Epoch 37/50
16/16 _____ 0s 24ms/step - loss: 6.6957e-04 - mae: 0.0201 - val_loss: 0.0038 - val_mae: 0.0483

Epoch 38/50
16/16 _____ 0s 22ms/step - loss: 0.0011 - mae: 0.0249 - val_loss: 0.0032 - val_mae: 0.0433

Epoch 39/50
16/16 _____ 1s 25ms/step - loss: 0.0014 - mae: 0.0280 - val_loss: 0.0033 - val_mae: 0.0442

Epoch 40/50
16/16 _____ 0s 27ms/step - loss: 9.4109e-04 - mae: 0.0234 - val_loss: 0.0031 - val_mae: 0.0425

Epoch 41/50
16/16 _____ 0s 25ms/step - loss: 7.6878e-04 - mae: 0.0206 - val_loss: 0.0035 - val_mae: 0.0463

Epoch 42/50
16/16 _____ 1s 24ms/step - loss: 0.0011 - mae: 0.0255 - val_loss: 0.0039 - val_mae: 0.0493

Epoch 43/50
16/16 _____ 0s 23ms/step - loss: 7.4514e-04 - mae: 0.0209 - val_loss: 0.0027 - val_mae: 0.0368

Epoch 44/50
16/16 _____ 0s 24ms/step - loss: 8.4085e-04 - mae: 0.0224 - val_loss: 0.0032 - val_mae: 0.0437

Epoch 45/50
16/16 _____ 0s 26ms/step - loss: 9.5099e-04 - mae: 0.0234 - val_loss: 0.0044 - val_mae: 0.0526

Epoch 46/50
16/16 _____ 0s 16ms/step - loss: 7.9463e-04 - mae: 0.0211 - val_loss: 0.0030 - val_mae: 0.0417

Epoch 47/50
16/16 _____ 0s 18ms/step - loss: 7.6148e-04 - mae: 0.0201 - val_loss: 0.0025 - val_mae: 0.0337


Epoch 48/50
16/16 _____ 0s 15ms/step - loss: 9.3954e-04 - mae: 0.0245 - val_loss: 0.0029 - val_mae: 0.0401

Epoch 49/50
16/16 _____ 0s 15ms/step - loss: 8.4489e-04 - mae: 0.0228 - val_loss: 0.0030 - val_mae: 0.0412

Epoch 50/50
16/16 _____ 0s 18ms/step - loss: 9.0218e-04 - mae: 0.0232 - val_loss: 0.0030 - val_mae: 0.0420

2/2 _____ 0s 263ms/step

RMSE: 28670.25
MAE : 21653.14



```
# 1. INSTALL REQUIRED LIBRARIES
!pip install -q pandas numpy scikit-learn matplotlib tensorflow

# 2. LOAD 'day_wise.csv'
import pandas as pd

df = pd.read_csv('/content/day_wise.csv')      # Adjust path if needed

# 3. PREPROCESSING
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error

df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)
df = df[['New cases']]

scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(df)

def create_sequences(data, seq_length=14):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i + seq_length])
        y.append(data[i + seq_length])
```

```

    return np.array(X), np.array(y)

seq_length = 14
X, y = create_sequences(scaled_data, seq_length)

split = int(len(X) * 0.8)
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]

X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))

# 4. BUILD AND TRAIN LSTM MODEL
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense

model = Sequential()
model.add(LSTM(50, activation='relu', input_shape=(seq_length, 1)))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error',
metrics=['mae'])

history = model.fit(X_train, y_train, epochs=50, batch_size=8,
validation_split=0.1, verbose=1)

# 5. PREDICT AND EVALUATE
y_pred = model.predict(X_test)
y_pred_inv = scaler.inverse_transform(y_pred)
y_test_inv = scaler.inverse_transform(y_test)

rmse = np.sqrt(mean_squared_error(y_test_inv, y_pred_inv))
mae = mean_absolute_error(y_test_inv, y_pred_inv)

print(f"RMSE: {rmse:.2f}")
print(f"MAE : {mae:.2f}")

# 6. PLOT PREDICTION VS ACTUAL
plt.figure(figsize=(12,6))
plt.plot(y_test_inv, label='Actual', marker='o')
plt.plot(y_pred_inv, label='Predicted', marker='x')
plt.title('Prediction vs Actual - New COVID-19 Cases')
plt.xlabel('Days')
plt.ylabel('New Cases')
plt.legend()
plt.grid(True)
plt.show()

```

Epoch 1/50

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/
rnn.py:200: UserWarning: Do not pass an `input_shape` / `input_dim`

```

argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

```
16/16 _____ 4s 53ms/step - loss: 0.0415 - mae: 0.1620 -  
val_loss: 0.0438 - val_mae: 0.2035  
Epoch 2/50  
16/16 _____ 1s 21ms/step - loss: 0.0077 - mae: 0.0749 -  
val_loss: 0.0079 - val_mae: 0.0745  
Epoch 3/50  
16/16 _____ 0s 13ms/step - loss: 0.0042 - mae: 0.0518 -  
val_loss: 0.0057 - val_mae: 0.0611  
Epoch 4/50  
16/16 _____ 0s 15ms/step - loss: 0.0025 - mae: 0.0414 -  
val_loss: 0.0028 - val_mae: 0.0371  
Epoch 5/50  
16/16 _____ 0s 13ms/step - loss: 0.0018 - mae: 0.0344 -  
val_loss: 0.0030 - val_mae: 0.0402  
Epoch 6/50  
16/16 _____ 0s 15ms/step - loss: 0.0015 - mae: 0.0314 -  
val_loss: 0.0026 - val_mae: 0.0345  
Epoch 7/50  
16/16 _____ 0s 15ms/step - loss: 0.0017 - mae: 0.0345 -  
val_loss: 0.0024 - val_mae: 0.0336  
Epoch 8/50  
16/16 _____ 0s 13ms/step - loss: 0.0015 - mae: 0.0302 -  
val_loss: 0.0026 - val_mae: 0.0360  
Epoch 9/50  
16/16 _____ 0s 13ms/step - loss: 0.0017 - mae: 0.0328 -  
val_loss: 0.0036 - val_mae: 0.0475  
Epoch 10/50  
16/16 _____ 0s 16ms/step - loss: 0.0015 - mae: 0.0307 -  
val_loss: 0.0025 - val_mae: 0.0335  
Epoch 11/50  
16/16 _____ 0s 13ms/step - loss: 0.0012 - mae: 0.0283 -  
val_loss: 0.0024 - val_mae: 0.0334  
Epoch 12/50  
16/16 _____ 0s 14ms/step - loss: 0.0015 - mae: 0.0317 -  
val_loss: 0.0029 - val_mae: 0.0393  
Epoch 13/50  
16/16 _____ 0s 13ms/step - loss: 0.0013 - mae: 0.0293 -  
val_loss: 0.0041 - val_mae: 0.0507  
Epoch 14/50  
16/16 _____ 0s 13ms/step - loss: 0.0015 - mae: 0.0308 -  
val_loss: 0.0025 - val_mae: 0.0337  
Epoch 15/50  
16/16 _____ 0s 15ms/step - loss: 0.0013 - mae: 0.0285 -  
val_loss: 0.0027 - val_mae: 0.0379  
Epoch 16/50  
16/16 _____ 0s 13ms/step - loss: 0.0010 - mae: 0.0260 -
```


val_loss: 0.0033 - val_mae: 0.0442
Epoch 17/50
16/16 _____ 0s 15ms/step - loss: 0.0012 - mae: 0.0280 -
val_loss: 0.0025 - val_mae: 0.0342
Epoch 18/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0260 -
val_loss: 0.0026 - val_mae: 0.0358
Epoch 19/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0269 -
val_loss: 0.0031 - val_mae: 0.0423
Epoch 20/50
16/16 _____ 0s 13ms/step - loss: 0.0010 - mae: 0.0255 -
val_loss: 0.0025 - val_mae: 0.0340
Epoch 21/50
16/16 _____ 0s 15ms/step - loss: 0.0012 - mae: 0.0277 -
val_loss: 0.0035 - val_mae: 0.0459
Epoch 22/50
16/16 _____ 0s 13ms/step - loss: 0.0012 - mae: 0.0305 -
val_loss: 0.0029 - val_mae: 0.0398
Epoch 23/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0257 -
val_loss: 0.0037 - val_mae: 0.0477
Epoch 24/50
16/16 _____ 0s 12ms/step - loss: 9.0265e-04 - mae:
0.0253 - val_loss: 0.0030 - val_mae: 0.0411
Epoch 25/50
16/16 _____ 0s 15ms/step - loss: 0.0010 - mae: 0.0246 -
val_loss: 0.0043 - val_mae: 0.0513
Epoch 26/50
16/16 _____ 0s 13ms/step - loss: 0.0012 - mae: 0.0254 -
val_loss: 0.0040 - val_mae: 0.0493
Epoch 27/50
16/16 _____ 0s 13ms/step - loss: 9.0281e-04 - mae:
0.0238 - val_loss: 0.0031 - val_mae: 0.0414
Epoch 28/50
16/16 _____ 0s 15ms/step - loss: 0.0012 - mae: 0.0263 -
val_loss: 0.0044 - val_mae: 0.0525
Epoch 29/50
16/16 _____ 0s 13ms/step - loss: 0.0010 - mae: 0.0251 -
val_loss: 0.0033 - val_mae: 0.0439
Epoch 30/50
16/16 _____ 0s 13ms/step - loss: 0.0013 - mae: 0.0276 -
val_loss: 0.0033 - val_mae: 0.0434
Epoch 31/50
16/16 _____ 0s 13ms/step - loss: 9.5889e-04 - mae:
0.0241 - val_loss: 0.0051 - val_mae: 0.0574
Epoch 32/50
16/16 _____ 0s 15ms/step - loss: 0.0011 - mae: 0.0258 -
val_loss: 0.0033 - val_mae: 0.0434

Epoch 33/50
16/16 _____ 0s 13ms/step - loss: 0.0010 - mae: 0.0248 -
val_loss: 0.0039 - val_mae: 0.0483

Epoch 34/50
16/16 _____ 0s 15ms/step - loss: 9.0899e-04 - mae:
0.0230 - val_loss: 0.0041 - val_mae: 0.0502

Epoch 35/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0248 -
val_loss: 0.0035 - val_mae: 0.0454

Epoch 36/50
16/16 _____ 0s 13ms/step - loss: 9.1811e-04 - mae:
0.0230 - val_loss: 0.0050 - val_mae: 0.0567

Epoch 37/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0267 -
val_loss: 0.0028 - val_mae: 0.0388

Epoch 38/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0258 -
val_loss: 0.0041 - val_mae: 0.0495

Epoch 39/50
16/16 _____ 0s 22ms/step - loss: 0.0011 - mae: 0.0268 -
val_loss: 0.0040 - val_mae: 0.0492

Epoch 40/50
16/16 _____ 0s 20ms/step - loss: 8.4081e-04 - mae:
0.0223 - val_loss: 0.0049 - val_mae: 0.0559

Epoch 41/50
16/16 _____ 0s 20ms/step - loss: 9.1679e-04 - mae:
0.0240 - val_loss: 0.0037 - val_mae: 0.0476

Epoch 42/50
16/16 _____ 1s 23ms/step - loss: 9.9494e-04 - mae:
0.0243 - val_loss: 0.0042 - val_mae: 0.0511

Epoch 43/50
16/16 _____ 0s 22ms/step - loss: 0.0010 - mae: 0.0240 -
val_loss: 0.0035 - val_mae: 0.0457

Epoch 44/50
16/16 _____ 0s 21ms/step - loss: 8.6277e-04 - mae:
0.0216 - val_loss: 0.0038 - val_mae: 0.0477

Epoch 45/50
16/16 _____ 1s 20ms/step - loss: 0.0011 - mae: 0.0260 -
val_loss: 0.0053 - val_mae: 0.0592

Epoch 46/50
16/16 _____ 0s 22ms/step - loss: 9.3478e-04 - mae:
0.0231 - val_loss: 0.0033 - val_mae: 0.0438

Epoch 47/50
16/16 _____ 1s 22ms/step - loss: 0.0011 - mae: 0.0247 -
val_loss: 0.0037 - val_mae: 0.0469


Epoch 48/50
16/16 _____ 1s 16ms/step - loss: 8.4657e-04 - mae:
0.0224 - val_loss: 0.0035 - val_mae: 0.0459

Epoch 49/50

```

16/16 _____ 0s 13ms/step - loss: 9.7624e-04 - mae:
0.0237 - val_loss: 0.0093 - val_mae: 0.0827
Epoch 50/50
16/16 _____ 0s 13ms/step - loss: 0.0011 - mae: 0.0252 -
val_loss: 0.0034 - val_mae: 0.0446
2/2 _____ 0s 213ms/step
RMSE: 28223.84
MAE : 24078.97

```



5.2. Sequence Text Prediction using LSTM

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from keras.utils import pad_sequences, to_categorical
from keras.models import Sequential
from keras.layers import Embedding, LSTM, Dense
import re

# Load Dataset
df = pd.read_csv('/content/sample_data/Jokes_subreddit_threads.csv')
# Adjust path as needed
df.dropna(inplace=True)

# Preprocess and Tokenize
text_data_list = df['selftext'].astype(str).tolist()           # Using
'selftext' for jokes

```

```

processed_text_data = [re.sub(r'^a-zA-Z\s.'], '', text) for text in
text_data_list]
text_data = ''.join(processed_text_data)
tokenizer = Tokenizer()
tokenizer.fit_on_texts(processed_text_data)          # Tokenize individual
jokes
total_words = len(tokenizer.word_index) + 1
print("Total words in vocab:", total_words)

# Create Sequences
input_sequences = []
tokens = tokenizer.texts_to_sequences([text_data])[0]

if len(tokens) > 1:
    for i in range(1, len(tokens)):
        n_gram_sequence = tokens[:i + 1]
        input_sequences.append(n_gram_sequence)

    max_seq_len = max([len(seq) for seq in input_sequences])
    input_sequences = pad_sequences(input_sequences,
maxlen=max_seq_len, padding='pre')

    X = input_sequences[:, :-1]
    y = input_sequences[:, -1]
    y = to_categorical(y, num_classes=total_words)

else:
    print("Not enough tokens to create sequences. Check your input
data or tokenization process.")
    X = np.array([])
    y = np.array([])
    max_seq_len = 0

# Build and Train LSTM Model
model = Sequential()
if max_seq_len > 0:
    model.add(Embedding(total_words, 100, input_length=max_seq_len -
1))
    model.add(LSTM(150))
    model.add(Dense(total_words, activation='softmax'))

    model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
    model.summary()

    history = model.fit(X, y, epochs=30, verbose=1)

# Plot Accuracy and Loss
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)

```

```

plt.plot(history.history['accuracy'], label='Accuracy')
plt.legend()
plt.title("Model Accuracy")

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Loss', color='red')
plt.legend()
plt.title("Model Loss")
plt.show()

# Text Generation Function
def generate_text(seed_text, next_words, model, max_seq_len):
    for _ in range(next_words):
        token_list = tokenizer.texts_to_sequences([seed_text])[0]
        token_list = pad_sequences([token_list],
maxlen=max_seq_len - 1, padding='pre')
        predicted = model.predict(token_list, verbose=0)
        predicted_word_index = np.argmax(predicted, axis=-1)[0]

        for word, index in tokenizer.word_index.items():
            if index == predicted_word_index:
                seed_text += " " + word
                break
    return seed_text

# Generate Sample Text
print("\nGenerated Text:\n")
print(generate_text("why did the chicken", 20, model,
max_seq_len))

else:
    print("Model cannot be built and trained due to insufficient
data.")

```

Total words in vocab: 28

```

<ipython-input-6-66999a8beb7e>:12: DtypeWarning: Columns (2) have
mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv('/content/sample_data/Jokes_subreddit_threads.csv')
# Adjust path as needed
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embeddin
g.py:90: UserWarning: Argument `input_length` is deprecated. Just
remove it.
warnings.warn(

```

Model: "sequential_1"

Layer (type)		Output Shape
Param #		

embedding (Embedding)	?	0
(unbuilt)		
lstm (LSTM)	?	0
(unbuilt)		
dense (Dense)	?	0
(unbuilt)		

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/30

1/1 ————— 3s 3s/step - accuracy: 0.0370 - loss: 3.3331

Epoch 2/30

1/1 ————— 0s 109ms/step - accuracy: 0.1481 - loss: 3.3245

Epoch 3/30

1/1 ————— 0s 78ms/step - accuracy: 0.1852 - loss: 3.3155

Epoch 4/30

1/1 ————— 0s 141ms/step - accuracy: 0.1481 - loss: 3.3055

Epoch 5/30

1/1 ————— 0s 97ms/step - accuracy: 0.1481 - loss: 3.2935

Epoch 6/30

1/1 ————— 0s 76ms/step - accuracy: 0.1481 - loss: 3.2781

Epoch 7/30

1/1 ————— 0s 142ms/step - accuracy: 0.1852 - loss: 3.2570

Epoch 8/30

1/1 ————— 0s 84ms/step - accuracy: 0.1481 - loss: 3.2260

Epoch 9/30

1/1 ————— 0s 80ms/step - accuracy: 0.1111 - loss: 3.1814

Epoch 10/30

1/1 ————— 0s 82ms/step - accuracy: 0.1111 - loss: 3.1360

Epoch 11/30
1/1 _____ 0s 77ms/step - accuracy: 0.1111 - loss:
3.1226

Epoch 12/30
1/1 _____ 0s 79ms/step - accuracy: 0.1111 - loss:
3.0747

Epoch 13/30
1/1 _____ 0s 140ms/step - accuracy: 0.1852 - loss:
3.0129

Epoch 14/30
1/1 _____ 0s 77ms/step - accuracy: 0.1852 - loss:
2.9704

Epoch 15/30
1/1 _____ 0s 160ms/step - accuracy: 0.1852 - loss:
2.9271

Epoch 16/30
1/1 _____ 0s 81ms/step - accuracy: 0.2222 - loss:
2.8563

Epoch 17/30
1/1 _____ 0s 142ms/step - accuracy: 0.1481 - loss:
2.7863

Epoch 18/30
1/1 _____ 0s 81ms/step - accuracy: 0.1481 - loss:
2.7374

Epoch 19/30
1/1 _____ 0s 79ms/step - accuracy: 0.1852 - loss:
2.6460

Epoch 20/30
1/1 _____ 0s 78ms/step - accuracy: 0.2963 - loss:
2.5975

Epoch 21/30
1/1 _____ 0s 79ms/step - accuracy: 0.3333 - loss:
2.5075

Epoch 22/30
1/1 _____ 0s 138ms/step - accuracy: 0.2963 - loss:
2.4693

Epoch 23/30
1/1 _____ 0s 141ms/step - accuracy: 0.2963 - loss:
2.3746

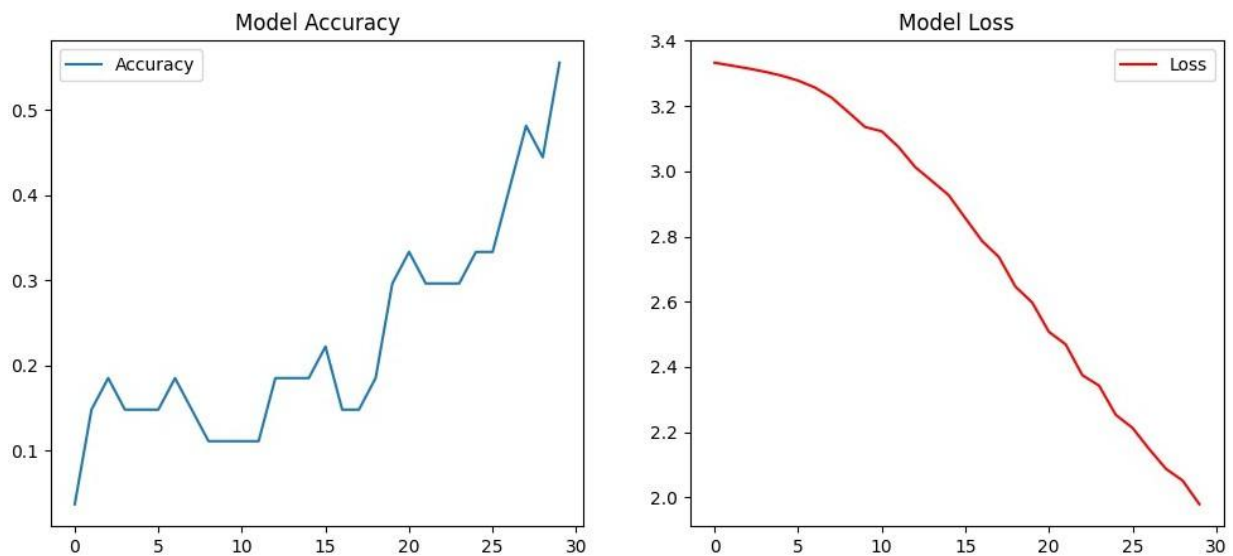
Epoch 24/30
1/1 _____ 0s 78ms/step - accuracy: 0.2963 - loss:
2.3427

Epoch 25/30
1/1 _____ 0s 99ms/step - accuracy: 0.3333 - loss:
2.2530

Epoch 26/30
1/1 _____ 0s 125ms/step - accuracy: 0.3333 - loss:
2.2130

Epoch 27/30

1/1 ————— 0s 77ms/step – accuracy: 0.4074 – loss: 2.1478
 Epoch 28/30
 1/1 ————— 0s 76ms/step – accuracy: 0.4815 – loss: 2.0875
 Epoch 29/30
 1/1 ————— 0s 78ms/step – accuracy: 0.4444 – loss: 2.0514
 Epoch 30/30
 1/1 ————— 0s 139ms/step – accuracy: 0.5556 – loss: 1.9788



Generated Text:

whv did the chicken has a a input input input punchline punchline your
 joke joke joke posting and if its been been been posted

5.3: Sequence Text Classification using LSTM

```
# Step 1: Install & Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
```



```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout

# Step 2: Load and Prepare Dataset
df = pd.read_csv("/content/sample_data/spam.csv", encoding='latin-1')
df = df[['v1', 'v2']]
df.columns = ['label', 'text']
df.dropna(inplace=True)

# Encode Labels: 'ham' = 0, 'spam' = 1
label_encoder = LabelEncoder()
df['label_num'] = label_encoder.fit_transform(df['label'])

# Check class distribution
print(df['label'].value_counts())

# Text Preprocessing (basic lowercase only)
df['text'] = df['text'].str.lower()

# Step 3: Tokenization and Padding
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['text'])
vocab_size = len(tokenizer.word_index) + 1

sequences = tokenizer.texts_to_sequences(df['text'])
max_len = max([len(seq) for seq in sequences]) # Maximum sequence
length
X = pad_sequences(sequences, maxlen=max_len)
y = df['label_num'].values

# Step 4: Split into Train/Test
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 5: Build and Compile LSTM Model
model = Sequential()
model.add(Embedding(input_dim=vocab_size, output_dim=128)) # Removed
deprecated input length
model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])

# Explicitly build the model to fix "unbuilt" summary issue
model.build(input_shape=(None, max_len))

# View model architecture
model.summary()

```

```

# Step 6: Train the Model
history = model.fit(X_train, y_train, epochs=5, batch_size=64,
validation_data=(X_test, y_test), verbose=1)

# Step 7: Evaluate Model
y_pred = (model.predict(X_test) > 0.5).astype("int32")

# Classification Report
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred, target_names=['ham',
'spam']))

# Accuracy Score
acc = accuracy_score(y_test, y_pred)
print("Accuracy:", acc)

# Step 8: Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Ham',
'Spam'], yticklabels=['Ham', 'Spam'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```

```

label
ham      4825
spam      747
Name: count, dtype: int64

```

```

Model: "sequential_3"

```

Layer (type)		Output Shape
Param #		
embedding_2 (Embedding)		(None, 189, 128)
1,141,888		
lstm_2 (LSTM)		(None, 128)
131,584		
dense_2 (Dense)		(None, 1)

129 |

Total params: 1,273,601 (4.86 MB)

Trainable params: 1,273,601 (4.86 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/5

70/70 ————— 47s 613ms/step - accuracy: 0.8886 - loss: 0.3466 - val_accuracy: 0.9794 - val_loss: 0.0804

Epoch 2/5

70/70 ————— 80s 582ms/step - accuracy: 0.9875 - loss: 0.0439 - val_accuracy: 0.9848 - val_loss: 0.0619

Epoch 3/5

70/70 ————— 41s 585ms/step - accuracy: 0.9946 - loss: 0.0195 - val_accuracy: 0.9830 - val_loss: 0.0750

Epoch 4/5

70/70 ————— 40s 570ms/step - accuracy: 0.9991 - loss: 0.0065 - val_accuracy: 0.9830 - val_loss: 0.0660

Epoch 5/5

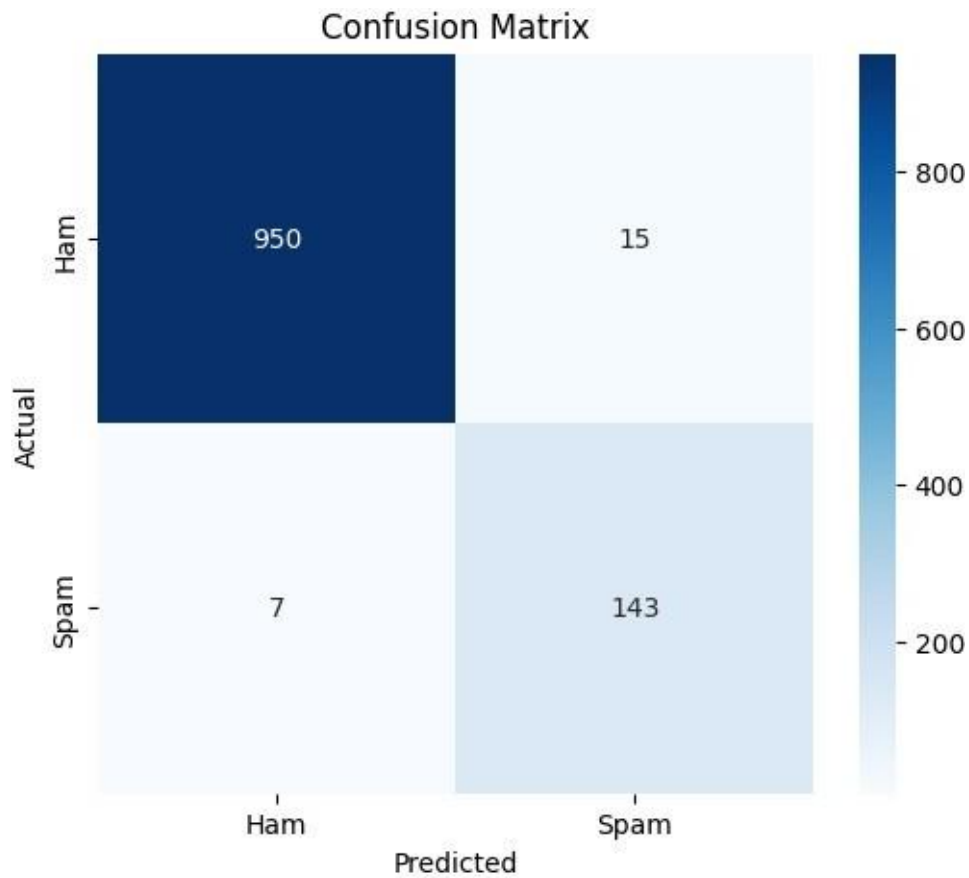
70/70 ————— 41s 568ms/step - accuracy: 0.9998 - loss: 0.0017 - val_accuracy: 0.9803 - val_loss: 0.0718

35/35 ————— 3s 76ms/step

Classification Report:

	precision	recall	f1-score	support
ham	0.99	0.98	0.99	965
spam	0.91	0.95	0.93	150
accuracy			0.98	1115
macro avg	0.95	0.97	0.96	1115
weighted avg	0.98	0.98	0.98	1115

Accuracy: 0.9802690582959641



Declaration

I, Manjiri Netankar, confirm that the work submitted in this assignment is my own and has been completed following academic integrity guidelines.

Github link:

Dataset Link : https://github.com/supriyamaskar/LSTM_deepLearning

Signature: Manjiri Netankar