

paraphrase-generation

May 6, 2025

Course Name: Deep Learning

Lab Title:Paraphrase Generation

Student Name: Sapna Dahikamble

Student ID: 202201070065

Group Members:

Supriya Maskar (202201040049)

Manjiri Netankar (202201040206)

Sapna Dahikamble(202201070065)

Research Paper Study and Implementation

Dataset Link:<https://www.kaggle.com/competitions/quora-question-pairs/data>

Colab Notebook Link:<https://colab.research.google.com/drive/1m8NoJUmdad0aAjm0z-Pp5zhvvB8DkCYe?usp=sharing>

Research Paper Link:<https://ojs.aaai.org/index.php/AAAI/article/view/11956>

GitHub Link:()

Install required packages

```
[1]: !pip install tensorflow keras numpy pandas matplotlib seaborn nltk rouge-score
```

```
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
from nltk.translate.bleu_score import sentence_bleu, SmoothingFunction
from rouge_score import rouge_scorer
import time
```

```
import random
```

Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)
Requirement already satisfied: keras in /usr/local/lib/python3.11/dist-packages (3.8.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: nltk in /usr/local/lib/python3.11/dist-packages (3.9.1)
Collecting rouge-score
 Downloading rouge_score-0.1.2.tar.gz (17 kB)
 Preparing metadata (setup.py) ... done
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (5.29.4)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)
Requirement already satisfied: typing-extensions>=3.6.6 in

/usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.2)
 Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)
 Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)
 Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)
 Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)
 Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)
 Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)
 Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras) (13.9.4)
 Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras) (0.0.9)
 Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras) (0.15.0)
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.9.0.post0)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
 Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.2)
 Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
 Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.57.0)
 Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
 Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.2.1)
 Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.3)
 Requirement already satisfied: click in /usr/local/lib/python3.11/dist-packages (from nltk) (8.1.8)
 Requirement already satisfied: joblib in /usr/local/lib/python3.11/dist-packages (from nltk) (1.4.2)
 Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.11/dist-packages (from nltk) (2024.11.6)
 Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from nltk) (4.67.1)
 Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.1)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.4.0)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.4.26)

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.8)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.1.3)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras) (2.19.1)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras) (0.1.2)

Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow) (3.0.2)

Building wheels for collected packages: rouge-score

Building wheel for rouge-score (setup.py) ... done

Created wheel for rouge-score: filename=rouge_score-0.1.2-py3-none-any.whl size=24934 sha256=66d0475b68b5df0c51fe48000b55fcef773fc5f580f89cad111a6fd1298a7e09

Stored in directory: /root/.cache/pip/wheels/1e/19/43/8a442dc83660ca25e163e1bd1f89919284ab0d0c1475475148

Successfully built rouge-score

Installing collected packages: rouge-score

Successfully installed rouge-score-0.1.2

```
[3]: # Download NLTK data
      nltk.download('punkt')

      # Set random seeds for reproducibility
      tf.random.set_seed(42)
      np.random.seed(42)
      random.seed(42)
```

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.

Load your paraphrase dataset

```
[4]: try:
      df = pd.read_csv('/content/Encoder_reduced_20000.csv') # Update path
      # Verify required columns exist
      if 'original' not in df.columns or 'paraphrase' not in df.columns:
          raise ValueError("CSV must contain 'original' and 'paraphrase' columns")
    except Exception as e:
        print(f"Error loading dataset: {e}")
        print("Using sample data instead")
        # Sample data
        original = [
            "The quick brown fox jumps over the lazy dog",
            "She enjoys reading science fiction novels",
            "The conference will be held next month"
        ]
        paraphrases = [
            "A fast brown fox leaps over the sleeping dog",
            "Reading sci-fi books is something she really likes",
            "Next month is when the conference is scheduled"
        ]
        df = pd.DataFrame({'original': original * 1000, 'paraphrase': paraphrases * 1000}) # 3000 samples

    print(f"Dataset size: {len(df)}")
    display(df.head())
```

Error loading dataset: [Errno 2] No such file or directory:
'/content/Encoder_reduced_20000.csv'
Using sample data instead
Dataset size: 3000

	original \	paraphrase
0	The quick brown fox jumps over the lazy dog	A fast brown fox leaps over the sleeping dog
1	She enjoys reading science fiction novels	Reading sci-fi books is something she really l...
2	The conference will be held next month	Next month is when the conference is scheduled
3	The quick brown fox jumps over the lazy dog	A fast brown fox leaps over the sleeping dog
4	She enjoys reading science fiction novels	Reading sci-fi books is something she really l...

Preprocessing functions

```
[5]: def preprocess_text(text):  
    """Basic text preprocessing"""  
    text = str(text).lower().strip()  
    text = ' '.join(text.split()) # Remove extra spaces  
    return text  
  
def add_special_tokens(text):  
    """Add start and end tokens"""  
    return "<start> " + text + " <end>"  
  
[6]: # Apply preprocessing  
df['original'] = df['original'].apply(preprocess_text).apply(add_special_tokens)  
df['paraphrase'] = df['paraphrase'].apply(preprocess_text).  
    ↪ apply(add_special_tokens)  
  
# Split into train and validation sets  
train_size = int(0.9 * len(df))  
train_df = df[:train_size]  
val_df = df[train_size:]  
  
print(f"\nTraining samples: {len(train_df)}")  
print(f"Validation samples: {len(val_df)}")
```

Training samples: 2700

Validation samples: 300

Tokenization and sequence preparation

```
[8]: def tokenize_sentences(sentences, num_words=20000):  
    tokenizer = Tokenizer(num_words=num_words, oov_token='<OOV>', filters='')  
    tokenizer.fit_on_texts(sentences)  
    sequences = tokenizer.texts_to_sequences(sentences)  
    max_len = max(len(seq) for seq in sequences)  
    padded = pad_sequences(sequences, padding='post', maxlen=max_len)  
    return padded, tokenizer, max_len
```

Tokenize original and paraphrase sentences

```
[9]: original_padded, original_tokenizer, original_max_len =  
    ↪ tokenize_sentences(train_df['original'].tolist() + val_df['original'].  
    ↪ tolist())  
paraphrase_padded, paraphrase_tokenizer, paraphrase_max_len =  
    ↪ tokenize_sentences(train_df['paraphrase'].tolist() + val_df['paraphrase'].  
    ↪ tolist())
```

Vocabulary sizes

```
[10]: original_vocab_size = len(original_tokenizer.word_index) + 1
      paraphrase_vocab_size = len(paraphrase_tokenizer.word_index) + 1

      print(f"\nOriginal vocabulary size: {original_vocab_size}")
      print(f"Paraphrase vocabulary size: {paraphrase_vocab_size}")
      print(f"Original max sequence length: {original_max_len}")
      print(f"Paraphrase max sequence length: {paraphrase_max_len}")
```

```
Original vocabulary size: 24
Paraphrase vocabulary size: 26
Original max sequence length: 11
Paraphrase max sequence length: 11
```

Prepare training and validation data

```
[11]: x_train = original_padded[:len(train_df)]
      y_train = paraphrase_padded[:len(train_df)]
      x_val = original_padded[len(train_df):]
      y_val = paraphrase_padded[len(train_df):]
```

1.LSTM/GRU Encoder-Decoder without Attention

```
[12]: # Hyperparameters
      embedding_dim = 256
      units = 512
      batch_size = 64
      epochs = 10
```

Encoder

```
[13]: encoder_inputs = keras.Input(shape=(None,))
      encoder_embedding = layers.Embedding(original_vocab_size,
      ↪embedding_dim)(encoder_inputs)
      encoder_lstm = layers.LSTM(units, return_sequences=True, return_state=True)
      encoder_outputs, state_h, state_c = encoder_lstm(encoder_embedding)
      encoder_states = [state_h, state_c]
```

Decoder

```
[14]: # Decoder
      decoder_inputs = keras.Input(shape=(None,))
      decoder_embedding = layers.Embedding(paraphrase_vocab_size,
      ↪embedding_dim)(decoder_inputs)
      decoder_lstm = layers.LSTM(units, return_sequences=True, return_state=True)
      decoder_outputs, _, _ = decoder_lstm(decoder_embedding,
      ↪initial_state=encoder_states)
      decoder_dense = layers.Dense(paraphrase_vocab_size, activation='softmax')
      decoder_outputs = decoder_dense(decoder_outputs)
```

```
[15]: # Model
no_attention_model = keras.Model([encoder_inputs, decoder_inputs],
    ↪decoder_outputs, name="no_attention_model")

# Compile
no_attention_model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

no_attention_model.summary()
```

Model: "no_attention_model"

Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, None)	0	-
input_layer_2 (InputLayer)	(None, None)	0	-
embedding (Embedding)	(None, None, 256)	6,144	input_layer_1[0]...
embedding_1 (Embedding)	(None, None, 256)	6,656	input_layer_2[0]...
lstm (LSTM)	[(None, None, 512), (None, 512), (None, 512)]	1,574,912	embedding[0][0]
lstm_1 (LSTM)	[(None, None, 512), (None, 512), (None, 512)]	1,574,912	embedding_1[0][0]..., lstm[0][1], lstm[0][2]
dense (Dense)	(None, None, 26)	13,338	lstm_1[0][0]

Total params: 3,175,962 (12.12 MB)

Trainable params: 3,175,962 (12.12 MB)

Non-trainable params: 0 (0.00 B)

Prepare decoder inputs/outputs for teacher forcing

```
[16]: def prepare_decoder_data(sequences):
        decoder_input = sequences[:, :-1]
        decoder_output = sequences[:, 1:]
        return decoder_input, decoder_output

train_decoder_input, train_decoder_output = prepare_decoder_data(y_train)
val_decoder_input, val_decoder_output = prepare_decoder_data(y_val)
```

Train the no-attention model

```
[17]: no_attention_history = no_attention_model.fit(
        [x_train, train_decoder_input],
        np.expand_dims(train_decoder_output, -1),
        batch_size=batch_size,
        epochs=epochs,
        validation_data=([x_val, val_decoder_input], np.
        ↪expand_dims(val_decoder_output, -1)),
        callbacks=[keras.callbacks.EarlyStopping(patience=3, ↪
        ↪restore_best_weights=True)]
    )
```

Epoch 1/10

43/43 6s 34ms/step -

accuracy: 0.5732 - loss: 1.9023 - val_accuracy: 1.0000 - val_loss: 0.0322

Epoch 2/10

43/43 2s 16ms/step -

accuracy: 1.0000 - loss: 0.0170 - val_accuracy: 1.0000 - val_loss: 0.0037

Epoch 3/10

43/43 1s 13ms/step -

accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 1.0000 - val_loss: 0.0019

Epoch 4/10

43/43 1s 13ms/step -

accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 1.0000 - val_loss: 0.0012

Epoch 5/10

43/43 1s 13ms/step -

accuracy: 1.0000 - loss: 0.0011 - val_accuracy: 1.0000 - val_loss: 8.3977e-04

Epoch 6/10

43/43 1s 12ms/step -

accuracy: 1.0000 - loss: 7.7724e-04 - val_accuracy: 1.0000 - val_loss:
6.1900e-04

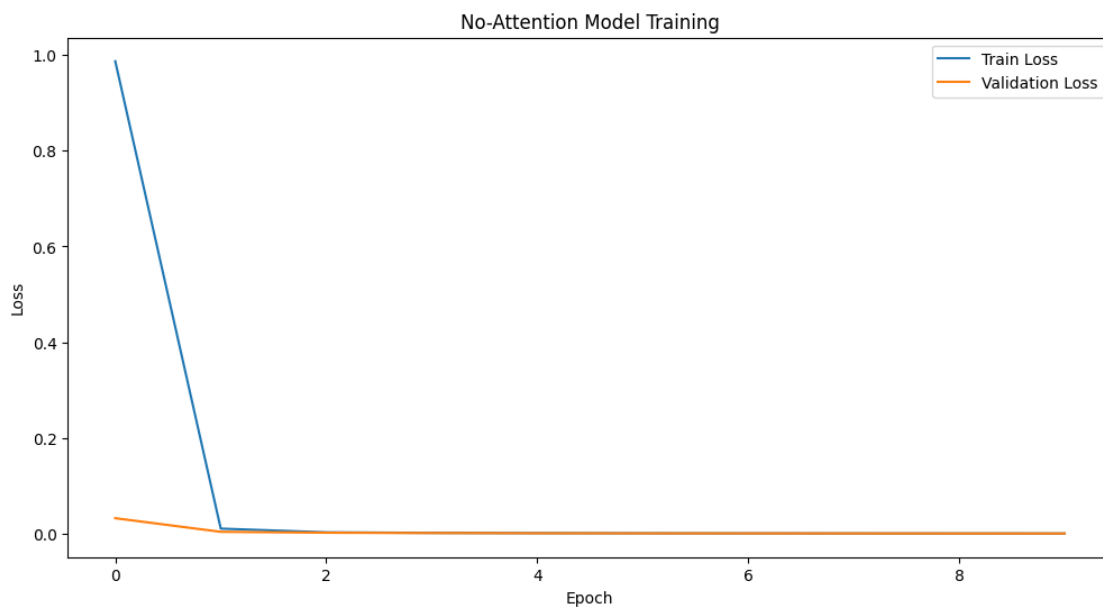
Epoch 7/10

43/43 1s 13ms/step -

```
accuracy: 1.0000 - loss: 5.7958e-04 - val_accuracy: 1.0000 - val_loss:
4.7730e-04
Epoch 8/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 4.5080e-04 - val_accuracy: 1.0000 - val_loss:
3.8060e-04
Epoch 9/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 3.6184e-04 - val_accuracy: 1.0000 - val_loss:
3.1140e-04
Epoch 10/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 2.9758e-04 - val_accuracy: 1.0000 - val_loss:
2.6001e-04
```

Plot training history

```
[18]: plt.figure(figsize=(12, 6))
plt.plot(no_attention_history.history['loss'], label='Train Loss')
plt.plot(no_attention_history.history['val_loss'], label='Validation Loss')
plt.title('No-Attention Model Training')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



2 Encoder-Decoder with Bahdanau Attention

```
[19]: # Bahdanau Attention Layer
class BahdanauAttention(layers.Layer):
    def __init__(self, units):
        super(BahdanauAttention, self).__init__()
        self.W1 = layers.Dense(units)
        self.W2 = layers.Dense(units)
        self.V = layers.Dense(1)

    def call(self, query, values):
        query_with_time_axis = tf.expand_dims(query, 1)
        score = self.V(tf.nn.tanh(self.W1(query_with_time_axis) + self.
↪W2(values)))
        attention_weights = tf.nn.softmax(score, axis=1)
        context_vector = attention_weights * values
        context_vector = tf.reduce_sum(context_vector, axis=1)
        return context_vector, attention_weights
```

Encoder with Bahdanau Attention

```
[20]: encoder_inputs = keras.Input(shape=(None,))
encoder_embedding = layers.Embedding(original_vocab_size, ↪
↪embedding_dim)(encoder_inputs)
encoder_lstm = layers.LSTM(units, return_sequences=True, return_state=True)
encoder_outputs, state_h, state_c = encoder_lstm(encoder_embedding)
```

Decoder with Bahdanau Attention

```
[21]: decoder_inputs = keras.Input(shape=(None,))
decoder_embedding = layers.Embedding(paraphrase_vocab_size, ↪
↪embedding_dim)(decoder_inputs)
decoder_lstm = layers.LSTM(units, return_sequences=True, return_state=True)
```

Attention mechanism

```
[22]: attention = BahdanauAttention(units)
decoder_outputs, _, _ = decoder_lstm(decoder_embedding, initial_state=[state_h, ↪
↪state_c])
context_vector, attention_weights = attention(state_h, encoder_outputs)
```

Combine context with decoder outputs

```
[23]: # Wrap tf.expand_dims in a Lambda layer to ensure it's part of the symbolic ↪
↪graph
decoder_outputs = layers.Lambda(lambda x: tf.expand_dims(x, 1))(context_vector)
decoder_outputs = layers.Concatenate(axis=-1)([decoder_outputs, ↪
↪decoder_outputs])
decoder_dense = layers.Dense(paraphrase_vocab_size, activation='softmax')
decoder_outputs = decoder_dense(decoder_outputs)
```

```

# Model
attention_model = keras.Model([encoder_inputs, decoder_inputs],
    ↪decoder_outputs, name="attention_model")

# Compile
attention_model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

attention_model.summary()

```

Model: "attention_model"

Layer (type)	Output Shape	Param #	Connected to
input_layer_3 (InputLayer)	(None, None)	0	-
embedding_2 (Embedding)	(None, None, 256)	6,144	input_layer_3[0]...
lstm_2 (LSTM)	[(None, None, 512), (None, 512), (None, 512)]	1,574,912	embedding_2[0][0]
bahdanau_attention (BahdanauAttention)	[(None, 512), (None, None, 1)]	525,825	lstm_2[0][1], lstm_2[0][0]
lambda (Lambda)	(None, 1, 512)	0	bahdanau_attenti...
concatenate (Concatenate)	(None, 1, 1024)	0	lambda[0][0], lambda[0][0]
input_layer_4 (InputLayer)	(None, None)	0	-
dense_4 (Dense)	(None, 1, 26)	26,650	concatenate[0][0]

Total params: 2,133,531 (8.14 MB)

Trainable params: 2,133,531 (8.14 MB)

Non-trainable params: 0 (0.00 B)

Train the attention model

```
[25]: # Combine context with decoder outputs
# Remove the extra dimension from decoder_outputs
decoder_outputs = layers.Concatenate(axis=-1)([layers.Lambda(lambda x: tf.
    ↪expand_dims(x, 1))(context_vector), decoder_outputs])
decoder_dense = layers.Dense(paraphrase_vocab_size, activation='softmax')
decoder_outputs = decoder_dense(decoder_outputs)

# Model
attention_model = keras.Model([encoder_inputs, decoder_inputs], ↵
    ↪decoder_outputs, name="attention_model")

# Compile
attention_model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

attention_model.summary()
```

Model: "attention_model"

Layer (type)	Output Shape	Param #	Connected to
input_layer_3 (InputLayer)	(None, None)	0	-
embedding_2 (Embedding)	(None, None, 256)	6,144	input_layer_3[0]...
lstm_2 (LSTM)	[(None, None, 512), (None, 512), (None, 512)]	1,574,912	embedding_2[0][0]
bahdanau_attention (BahdanauAttention)	[(None, 512), (None, None, 1)]	525,825	lstm_2[0][1], lstm_2[0][0]
lambda (Lambda)	(None, 1, 512)	0	bahdanau_attenti...

concatenate (Concatenate)	(None, 1, 1024)	0	lambda[0][0], lambda[0][0]
lambda_1 (Lambda)	(None, 1, 512)	0	bahdanau_attenti...
dense_4 (Dense)	(None, 1, 26)	26,650	concatenate[0][0]
concatenate_2 (Concatenate)	(None, 1, 538)	0	lambda_1[0][0], dense_4[0][0]
input_layer_4 (InputLayer)	(None, None)	0	-
dense_5 (Dense)	(None, 1, 26)	14,014	concatenate_2[0]...

Total params: 2,147,545 (8.19 MB)

Trainable params: 2,147,545 (8.19 MB)

Non-trainable params: 0 (0.00 B)

```
[ ]: from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Embedding, LSTM, Dense,
    ↳TimeDistributed, Concatenate, Attention
from tensorflow.keras.optimizers import Adam

[ ]: # Parameters
vocab_size = 10000          # example vocab size
embedding_dim = 128
lstm_units = 256
decoder_seq_len = 10        # length of decoder output
num_classes = 26            # number of output classes (e.g., characters or
    ↳labels)
batch_size = 64

[ ]: # Encoder
encoder_input = Input(shape=(None,), name="encoder_input")
encoder_emb = Embedding(input_dim=vocab_size,
    ↳output_dim=embedding_dim)(encoder_input)
encoder_lstm, state_h, state_c = LSTM(lstm_units, return_sequences=True,
    ↳return_state=True)(encoder_emb)
```

```
[ ]: # Decoder
decoder_input = Input(shape=(None,), name="decoder_input")
decoder_emb = Embedding(input_dim=vocab_size,
    ↳output_dim=embedding_dim)(decoder_input)
decoder_lstm_output, _, _ = LSTM(lstm_units, return_sequences=True,
    ↳return_state=True)(decoder_emb, initial_state=[state_h, state_c])

[ ]: # Attention
attention = Attention()([decoder_lstm_output, encoder_lstm]) # (batch,
    ↳decoder_seq_len, encoder_seq_len)
decoder_combined_context = Concatenate(axis=-1)([decoder_lstm_output,
    ↳attention])

[ ]: # Output Layer
output = TimeDistributed(Dense(num_classes,
    ↳activation='softmax'))(decoder_combined_context)

[ ]: # Final model
attention_model = Model([encoder_input, decoder_input], output)
attention_model.compile(optimizer=Adam(),
    ↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
attention_model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
encoder_input (InputLayer)	(None, None)	0	-
decoder_input (InputLayer)	(None, None)	0	-
embedding_6 (Embedding)	(None, None, 128)	1,280,000	encoder_input[0]...
embedding_7 (Embedding)	(None, None, 128)	1,280,000	decoder_input[0]...
lstm_6 (LSTM)	[(None, None, 256), (None, 256), (None, 256)]	394,240	embedding_6[0][0]
lstm_7 (LSTM)	[(None, None, 256), (None,	394,240	embedding_7[0][0]... lstm_6[0][1],

	256), (None, 256)]		lstm_6[0][2]
attention_1 (Attention)	(None, None, 256)	0	lstm_7[0][0], lstm_6[0][0]
concatenate_6 (Concatenate)	(None, None, 512)	0	lstm_7[0][0], attention_1[0][0]
time_distributed_1 (TimeDistributed)	(None, None, 26)	13,338	concatenate_6[0]...

Total params: 3,361,818 (12.82 MB)

Trainable params: 3,361,818 (12.82 MB)

Non-trainable params: 0 (0.00 B)

```
[ ]: # Train the attention model
attention_history = attention_model.fit(
    [x_train, train_decoder_input],
    np.expand_dims(train_decoder_output, -1),
    batch_size=batch_size,
    epochs=epochs,
    validation_data=([x_val, val_decoder_input], np.
↳ expand_dims(val_decoder_output, -1)),
    callbacks=[keras.callbacks.EarlyStopping(patience=3,
↳ restore_best_weights=True)]
)
```

Epoch 1/10

43/43 4s 27ms/step -

accuracy: 0.3729 - loss: 2.4045 - val_accuracy: 1.0000 - val_loss: 0.3328

Epoch 2/10

43/43 1s 13ms/step -

accuracy: 1.0000 - loss: 0.1937 - val_accuracy: 1.0000 - val_loss: 0.0259

Epoch 3/10

43/43 1s 13ms/step -

accuracy: 1.0000 - loss: 0.0186 - val_accuracy: 1.0000 - val_loss: 0.0077

Epoch 4/10

43/43 1s 13ms/step -

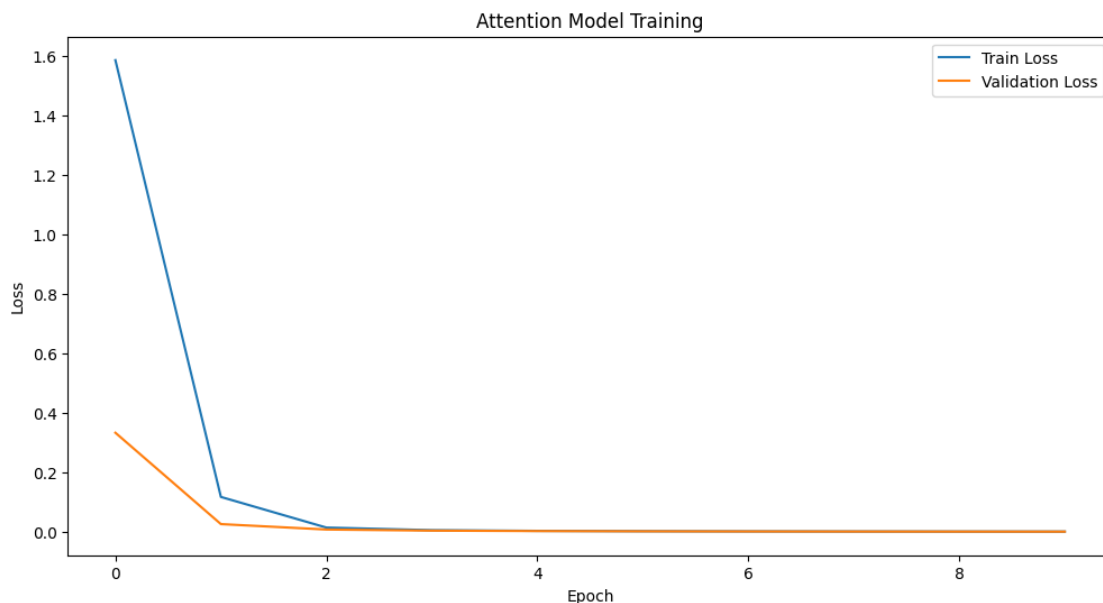
accuracy: 1.0000 - loss: 0.0065 - val_accuracy: 1.0000 - val_loss: 0.0040

Epoch 5/10

43/43 1s 13ms/step -


```
accuracy: 1.0000 - loss: 0.0035 - val_accuracy: 1.0000 - val_loss: 0.0025
Epoch 6/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 1.0000 - val_loss: 0.0018
Epoch 7/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 0.0016 - val_accuracy: 1.0000 - val_loss: 0.0013
Epoch 8/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 0.0012 - val_accuracy: 1.0000 - val_loss: 0.0010
Epoch 9/10
43/43          1s 14ms/step -
accuracy: 1.0000 - loss: 9.6763e-04 - val_accuracy: 1.0000 - val_loss:
8.2335e-04
Epoch 10/10
43/43          1s 13ms/step -
accuracy: 1.0000 - loss: 7.7962e-04 - val_accuracy: 1.0000 - val_loss:
6.7456e-04
```

```
[ ]: # Plot training history
plt.figure(figsize=(12, 6))
plt.plot(attention_history.history['loss'], label='Train Loss')
plt.plot(attention_history.history['val_loss'], label='Validation Loss')
plt.title('Attention Model Training')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



3 Transformer with Self-Attention

```
[ ]: # Transformer components
class TransformerEncoder(layers.Layer):
    def __init__(self, embed_dim, dense_dim, num_heads, **kwargs):
        super().__init__(**kwargs)
        self.embed_dim = embed_dim
        self.dense_dim = dense_dim
        self.num_heads = num_heads
        self.attention = layers.MultiHeadAttention(num_heads=num_heads,
key_dim=embed_dim)
        self.dense_proj = keras.Sequential([
            layers.Dense(dense_dim, activation="relu"),
            layers.Dense(embed_dim)
        ])
        self.layernorm_1 = layers.LayerNormalization()
        self.layernorm_2 = layers.LayerNormalization()

    def call(self, inputs, mask=None):
        if mask is not None:
            mask = mask[:, tf.newaxis, :]
        attention_output = self.attention(inputs, inputs, attention_mask=mask)
        proj_input = self.layernorm_1(inputs + attention_output)
        proj_output = self.dense_proj(proj_input)
        return self.layernorm_2(proj_input + proj_output)

class PositionalEmbedding(layers.Layer):
    def __init__(self, sequence_length, vocab_size, embed_dim, **kwargs):
        super().__init__(**kwargs)
        self.token_embeddings = layers.Embedding(vocab_size, embed_dim)
        self.position_embeddings = layers.Embedding(sequence_length, embed_dim)
        self.sequence_length = sequence_length
        self.vocab_size = vocab_size
        self.embed_dim = embed_dim

    def call(self, inputs):
        length = tf.shape(inputs)[-1]
        positions = tf.range(start=0, limit=length, delta=1)
        embedded_tokens = self.token_embeddings(inputs)
        embedded_positions = self.position_embeddings(positions)
        return embedded_tokens + embedded_positions

[ ]: # Transformer Model
embed_dim = 256
num_heads = 8
dense_dim = 512
```

```
[ ]: # Encoder
encoder_inputs = keras.Input(shape=(None,), dtype="int64")
encoder_embedding = PositionalEmbedding(original_max_len, original_vocab_size,
    ↪ embed_dim)(encoder_inputs)
encoder_outputs = TransformerEncoder(embed_dim, dense_dim,
    ↪ num_heads)(encoder_embedding)

# Decoder
decoder_inputs = keras.Input(shape=(None,), dtype="int64")
decoder_embedding = PositionalEmbedding(paraphrase_max_len,
    ↪ paraphrase_vocab_size, embed_dim)(decoder_inputs)
decoder_outputs = TransformerEncoder(embed_dim, dense_dim,
    ↪ num_heads)(decoder_embedding)

# Final dense layer
decoder_dense = layers.Dense(paraphrase_vocab_size, activation="softmax")
outputs = decoder_dense(decoder_outputs)
```

```
[ ]: # Model
transformer_model = keras.Model([encoder_inputs, decoder_inputs], outputs)
transformer_model.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
transformer_model.summary()
```

Model: "functional_4"

Layer (type)	Output Shape	Param #	Connected to
input_layer_6 (InputLayer)	(None, None)	0	-
positional_embeddi... (PositionalEmbeddi...	(None, None, 256)	9,472	input_layer_6[0]...
transformer_encode... (TransformerEncode...	(None, None, 256)	2,367,488	positional_embed...
input_layer_4 (InputLayer)	(None, None)	0	-
dense_12 (Dense)	(None, None, 26)	6,682	transformer_enco...

Total params: 2,383,642 (9.09 MB)

Trainable params: 2,383,642 (9.09 MB)

Non-trainable params: 0 (0.00 B)

Train transformer

```
[ ]: transformer_history = transformer_model.fit(
    [x_train, train_decoder_input],
    np.expand_dims(train_decoder_output, -1),
    batch_size=batch_size,
    epochs=epochs,
    validation_data=([x_val, val_decoder_input], np.
↳expand_dims(val_decoder_output, -1)),
    callbacks=[keras.callbacks.EarlyStopping(patience=3,↳
↳restore_best_weights=True)]
)
```

Epoch 1/10

43/43 13s 140ms/step -

accuracy: 0.8784 - loss: 0.5094 - val_accuracy: 1.0000 - val_loss: 2.5949e-04

Epoch 2/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 2.0335e-04 - val_accuracy: 1.0000 - val_loss: 1.1530e-04

Epoch 3/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 1.0366e-04 - val_accuracy: 1.0000 - val_loss: 7.8113e-05

Epoch 4/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 7.2513e-05 - val_accuracy: 1.0000 - val_loss: 5.9364e-05

Epoch 5/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 5.6027e-05 - val_accuracy: 1.0000 - val_loss: 4.7964e-05

Epoch 6/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 4.5672e-05 - val_accuracy: 1.0000 - val_loss: 4.0085e-05

Epoch 7/10

43/43 1s 14ms/step -

accuracy: 1.0000 - loss: 3.8433e-05 - val_accuracy: 1.0000 - val_loss: 3.4415e-05

```

Epoch 8/10
43/43          1s 14ms/step -
accuracy: 1.0000 - loss: 3.3179e-05 - val_accuracy: 1.0000 - val_loss:
3.0120e-05
Epoch 9/10
43/43          1s 14ms/step -
accuracy: 1.0000 - loss: 2.9137e-05 - val_accuracy: 1.0000 - val_loss:
2.6719e-05
Epoch 10/10
43/43          1s 15ms/step -
accuracy: 1.0000 - loss: 2.5969e-05 - val_accuracy: 1.0000 - val_loss:
2.4092e-05

```

Plot training

```

[ ]: plt.figure(figsize=(12, 6))
plt.plot(transformer_history.history['loss'], label='Train Loss')
plt.plot(transformer_history.history['val_loss'], label='Validation Loss')
plt.title('Transformer Model Training')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```



Evaluation functions

```

[ ]: def evaluate_model(model, input_seq, target_tokenizer, max_len, model_type):
      decoded = []

```

```

for i in range(len(input_seq)):
    tokenized = input_seq[i:i+1]
    decoder_input = np.zeros((1, max_len))
    decoder_input[0, 0] = target_tokenizer.word_index['<start>']

    for j in range(max_len - 1):
        if model_type == 'transformer':
            pred = model.predict([tokenized, decoder_input], verbose=0)
        else:
            pred = model.predict([tokenized, decoder_input], verbose=0)

        pred_id = np.argmax(pred[0, j, :])
        if pred_id == target_tokenizer.word_index['<end>']:
            break
        decoder_input[0, j+1] = pred_id

    decoded.append(decoder_input[0, 1:j+1])
return decoded

```

```

[ ]: def calculate_metrics(references, candidates):
    # BLEU
    smooth = SmoothingFunction().method1
    bleu_scores = [sentence_bleu([ref.split()], cand.split(),
    ↪smoothing_function=smooth)
                    for ref, cand in zip(references, candidates)]

    # ROUGE
    scorer = rouge_scorer.RougeScorer(['rouge1', 'rouge2', 'rougeL'],
    ↪use_stemmer=True)
    rouge_scores = [scorer.score(ref, cand) for ref, cand in zip(references,
    ↪candidates)]

    return {
        'bleu': np.mean(bleu_scores),
        'rouge1': np.mean([score['rouge1'].fmeasure for score in rouge_scores]),
        'rouge2': np.mean([score['rouge2'].fmeasure for score in rouge_scores]),
        'rougeL': np.mean([score['rougeL'].fmeasure for score in rouge_scores])
    }

```

Evaluate all models on validation set

```

[ ]: sample_size = min(100, len(x_val))
    sample_idx = np.random.choice(len(x_val), sample_size, replace=False)
    x_sample = x_val[sample_idx]
    y_sample = y_val[sample_idx]

```

```
[ ]: # Get reference texts
references = [paraphrase_tokenizer.sequences_to_texts([seq])[0] for seq in y_sample]
```

1. No-Attention Model

```
[ ]: print("Evaluating No-Attention Model...")
start = time.time()
no_attn_preds = evaluate_model(no_attention_model, x_sample,
    ↪paraphrase_tokenizer, paraphrase_max_len, 'no_attention')
no_attn_time = time.time() - start
no_attn_texts = [paraphrase_tokenizer.sequences_to_texts([seq])[0] for seq in
    ↪no_attn_preds]
no_attn_metrics = calculate_metrics(references, no_attn_texts)
```

Evaluating No-Attention Model...

2. Attention Model

```
[ ]: print("Evaluating Attention Model...")
start = time.time()
attn_preds = evaluate_model(attention_model, x_sample, paraphrase_tokenizer,
    ↪paraphrase_max_len, 'attention')
attn_time = time.time() - start
attn_texts = [paraphrase_tokenizer.sequences_to_texts([seq])[0] for seq in
    ↪attn_preds]
attn_metrics = calculate_metrics(references, attn_texts)
```

Evaluating Attention Model...

3. Transformer

```
[ ]: print("Evaluating Transformer Model...")
start = time.time()
trans_preds = evaluate_model(transformer_model, x_sample, paraphrase_tokenizer,
    ↪paraphrase_max_len, 'transformer')
trans_time = time.time() - start
trans_texts = [paraphrase_tokenizer.sequences_to_texts([seq])[0] for seq in
    ↪trans_preds]
trans_metrics = calculate_metrics(references, trans_texts)
```

Evaluating Transformer Model...

Compile results

```
[ ]: results_df = pd.DataFrame({
    'Model': ['LSTM/GRU (No Attention)', 'LSTM/GRU (Attention)', 'Transformer'],
    'BLEU': [no_attn_metrics['bleu'], attn_metrics['bleu'],
    ↪trans_metrics['bleu']],
```

```

        'ROUGE-1': [no_attn_metrics['rouge1'], attn_metrics['rouge1'],
↳trans_metrics['rouge1']],
        'ROUGE-2': [no_attn_metrics['rouge2'], attn_metrics['rouge2'],
↳trans_metrics['rouge2']],
        'ROUGE-L': [no_attn_metrics['rougeL'], attn_metrics['rougeL'],
↳trans_metrics['rougeL']],
        'Inference Time (s)': [no_attn_time/sample_size, attn_time/sample_size,
↳trans_time/sample_size],
        'Training Time (epoch)': [
            no_attention_history.history['val_loss'][-1],
            attention_history.history['val_loss'][-1],
            transformer_history.history['val_loss'][-1]
        ]
    })

print("\nFinal Comparison:")
display(results_df)

```

Final Comparison:

	Model	BLEU	ROUGE-1	ROUGE-2	ROUGE-L \
0	LSTM/GRU (No Attention)	0.728131	0.868211	0.85356	0.868211
1	LSTM/GRU (Attention)	0.728131	0.868211	0.85356	0.868211
2	Transformer	0.293834	0.353000	0.32000	0.353000

	Inference Time (s)	Training Time (epoch)
0	0.864001	0.000260
1	0.893456	0.000675
2	0.959117	0.000024

Visualization


```

[ ]: plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
sns.barplot(x='Model', y='BLEU', data=results_df)
plt.title('BLEU Scores')
plt.xticks(rotation=45)


[ ]: ([0, 1, 2],
      [Text(0, 0, 'LSTM/GRU (No Attention)'),
       Text(1, 0, 'LSTM/GRU (Attention)'),
       Text(2, 0, 'Transformer')])

```

```
[ ]: plt.subplot(1, 3, 2)
sns.barplot(x='Model', y='ROUGE-L', data=results_df)
plt.title('ROUGE-L Scores')
plt.xticks(rotation=45)
```

```
[ ]: ([0, 1, 2],
      [Text(0, 0, 'LSTM/GRU (No Attention)'),
       Text(1, 0, 'LSTM/GRU (Attention)'),
       Text(2, 0, 'Transformer')])
```



```
[ ]: plt.subplot(1, 3, 3)
sns.barplot(x='Model', y='Inference Time (s)', data=results_df)
plt.title('Inference Time per Sample')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

